

Using the data.table package

Steve Pittard wsp@emory.edu

November 6, 2015

Motivations

- A data frame is a set of columns. Every column is same length but of possibly different types.
- It has characteristics of both a matrix, (each row is the same data type),
- Each column can be a different data type
- Bracket notation offers a convenient way to search through the data drame

Motivations

```
DT <- data.frame(x=c("B","A","B","A","B"),y=c(7,2,1,5,9))
```

```
  x y
1 B 7
2 A 2
3 B 1
4 A 5
5 B 9
```

```
DT[2:3,]
```

```
  x y
2 A 2
3 B 1
```

Get Rows 2 and 3

DT[2:3,]

X	Y
B	7
A	2
B	1
A	5
B	9

Find all rows where x is equal to B and y greater than 2

```
DT[DT$x=="B" & DT$y > 2,]
```

X	Y
B	7
A	2
B	1
A	5
B	9

Find all rows where x is equal to B and y greater than 2

If DT were a table in some relational database system such as MySQL, SQLite, Oracle, Postgress, Access, etc then you could use SQL to extract the information. So if you know SQL then this approach might be more attractive to you.

```
select * from DT where x = 'B' and y > 2
```

X	Y
B	7
A	2
B	1
A	5
B	9

Find all rows where x is equal to B and y greater than 2

Note that there is an R package called **sqldf** that lets you treat a data frame as if it were a table in a relational database.

```
library(sqldf)  
sqldf("select * from DT where x = 'B' and y > 2")
```

```
  x y  
1 B 7  
2 B 9
```

X	Y
B	7
A	2
B	1
A	5
B	9

Motivations

But some things that are natural to want to do do not work with standard data frames

```
DT[, 'y']  
[1] 7 2 1 5 9
```

I think the following should work but it does not

```
DT[, sum('y')]  
Error in sum("y") : invalid 'type' (character) of argument
```

```
sum(DT$y)           # You have to do it this way  
[1] 24
```


Summarizing by Factors

We frequently want to summarize by some numeric value in terms of factors. For example find the average value of y for each level of factor x . R has ways to do this. There are functions we can call

```
tapply(DT$y, DT$x, mean)
```

	A	B
	3.500000	5.666667

OR

```
aggregate(y~., data=DT, mean)
```

	x	y
1	A	3.500000
2	B	5.666667

Summarizing by Factors

But wouldn't it be nice to have summary capability as part of the data frame structure ?

data.table allows us to do just that. There are packages such as dplyr that do this als but we will look at **data.table** first.

Some advantages of **data.table** is that:

- Works well with huge data files
- Extends the data frame bracket notation to do more
- Works with non data.table aware functions (that is it can act just like a regular data frame)

Summarizing by Factors

We will turn DT into a data.table - it is easy and it will still act like a data frame if you use the standard data frame functions

```
NDT <- data.table(DT)
```

```
class(NDT)
```

```
[1] "data.table" "data.frame"
```

```
nrow(NDT)
```

```
[1] 5
```

```
NDT[,sum(y)] # Aha - this will not work with typical data frames
```

```
[1] 24
```

Summarizing by Factors

Aggregating with **data.table** follows a pattern very similar to SQL although you don't need to know SQL to memorize it

DT[i, j, by]

Reading the above out loud would sound like: "With DT, subset rows using i, then calculate j as grouped by k."

R	:	i	j	by
SQL	:	WHERE	SELECT	GROUP BY

Summarizing by Factors

Using NDT, for all rows calculate the sum and mean of y

```
NDT[,.(total = sum(y), mean = mean(y))]
```

```
      total mean
```

```
1:      24  4.8
```

For all rows calculate sum and mean of y and group by x

```
NDT[,.(total = sum(y), mean = mean(y)),by=x]
```

```
      x total      mean
```

```
1: B      17 5.666667
```

```
2: A       7 3.500000
```

Old way - you have to use a function like aggregate

```
aggregate(y~x,data=NDT,function(v) c(total=sum(v),mean=mean(v)))
```

```
      x  y.total  y.mean
```

```
1 A  7.000000  3.500000
```

```
2 B 17.000000  5.666667
```

Summarizing by Factors

How many rows in NDT ?

```
NDT[,.N]  
[1] 5
```

How many observations in each group

```
NDT[,.N,by=x]  
  x N  
1: B 3  
2: A 2
```

For rows where $y > 2$ how many observations in each group

```
NDT[y > 2, .N, by=x]  
  x N  
1: B 2  
2: A 1
```

Summarizing by Factors

Let's do some work with the built in mtcars data frame

```
tcars <- data.table(mtcars)
```

```
str(tcars)
```

```
Classes data.table and data.frame: 32 obs. of 11 variables:
```

```
$ mpg : num 21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
```

```
$ cyl : num 6 6 4 6 8 6 8 4 4 6 ...
```

```
$ disp: num 160 160 108 258 360 ...
```

```
$ hp : num 110 110 93 110 175 105 245 62 95 123 ...
```

```
$ drat: num 3.9 3.9 3.85 3.08 3.15 2.76 3.21 3.69 3.92 3.92 ...
```

```
$ wt : num 2.62 2.88 2.32 3.21 3.44 ...
```

```
$ qsec: num 16.5 17 18.6 19.4 17 ...
```

```
$ vs : num 0 0 1 1 0 1 0 1 1 1 ...
```

```
$ am : num 1 1 1 0 0 0 0 0 0 0 ...
```

```
$ gear: num 4 4 4 3 3 3 3 4 4 4 ...
```

```
$ carb: num 4 4 1 1 2 1 4 2 2 4 ...
```

```
- attr(*, ".internal.selfref")=<externalptr>
```

Summarizing by Factors

For all rows calculate the mean mpg as grouped by trans type

```
tcars[,.(mean_mpg = mean(mpg)),by=.(trans=am)]
```

	trans	mean_mpg
1:	1	24.39231
2:	0	17.14737

For all rows calculate mean mpg as grouped by trans type and cylinder

```
tcars[,.(mean_mpg = mean(mpg)),by=.(trans=am,cyl=cyl)]
```

	trans	cyl	mean_mpg
1:	1	6	20.56667
2:	1	4	28.07500
3:	0	6	19.12500
4:	0	8	15.05000
5:	0	4	22.90000
6:	1	8	15.40000

Summarizing by Factors

```
# To order the result in descending order of mean_mpg we just treat it like  
# data frame - which it can act like !
```

```
tcars[,.(mean_mpg = mean(mpg)),by=.(trans=am,cyl=cyl)][order(-mean_mpg),]
```

	trans	cyl	mean_mpg
1:	1	4	28.07500
2:	0	4	22.90000
3:	1	6	20.56667
4:	0	6	19.12500
5:	1	8	15.40000
6:	0	8	15.05000

```
# Could have done
```

```
tmpdf <- tcars[,.(mean_mpg = mean(mpg)),by=.(trans=am,cyl=cyl)]  
tmpdf[order(-tmpdf$mean_mpg),]
```

Summarizing by Factors

where car weight > 2.5 tons calculate the mean mpg as grouped by gears

```
tcars[wt > 2.5,.(mean_mpg = mean(mpg)), by=.(gear=gear)]
```

```
  gear mean_mpg
1:    4 21.08571
2:    3 15.72143
3:    5 16.83333
```

```
tcars[,.N] # How many rows ?
```

```
[1] 32
```

```
tcars[,.N,by=am] # Number of obs per trans type
```

```
  am  N
1:  1 13
2:  0 19
```

```
tcars[,.N,by=cyl] # Number of obs per cyl group
```

```
  cyl  N
1:   6   7
2:   4  11
3:   8  14
```

Example: Reading in a “large” .csv file

- Consider a .csv file that has 334,142 rows relating to calls to the Chicago police in 2012
- http://stevie42.bitbucket.org/YOUTUBE.DIR/chi_crimes.csv
- Use the default **read.csv** function to create a data frame
- Let's time this function call to see how long it will take on my MacBook Pro with 8GB of RAM:

```
system.time(df.crimes <- read.csv("chi_crimes.csv",  
                                header=TRUE, sep=","))
```

```
   user  system elapsed  
30.251   0.283   30.569
```

```
nrow(df.crimes)  
[1] 334141
```

Example: Reading in a “large” .csv file

Let's read in this file using **data.table** function. The **fread** function returns a data table

```
library(data.table)
```

```
system.time(dt.crimes <- fread("chi_crimes.csv",  
                               header=TRUE, sep=","))
```

```
   user  system elapsed  
1.045   0.037   1.362
```

```
attributes(dt.crimes)$class      # dt.crimes is also a data.frame  
[1] "data.table" "data.frame"
```

```
nrow(df.crimes)  
[1] 334141
```

```
dt.crimes[, .N]  
[1] 334141
```

Aggregation on a large file

This data frame has information on every call to Chicago police in the year 2013.

So we'll want to see what factors there are in the dataframe so we can do some summaries across groups.

```
names(df.crimes)
```

[1] "Case.Number"	"ID"	"Date"
[5] "IUCR"	"Primary.Type"	"Description"
[9] "Arrest"	"Domestic"	"Beat"
[13] "Ward"	"FBI.Code"	"X.Coordinate"
[17] "Y.Coordinate"	"Year"	"Latitude"

Aggregation on a large file

Let's see how many unique values there are for each column. Looks like 30 FBI codes so maybe we could see the number of calls per FBI code. What about District ? There are 25 of those.

```
sapply(df.crimes,function(x) {length(unique(x))})
```

Case.Number	ID	Date
334114	334139	121484
Primary.Type	Description	Location.Description
30	296	120
Beat	District	Ward
302	25	51
Community.Area	Y.Coordinate	Year
79	89895	1
Longitude	Location	
180393	178534	

Aggregation on a Large File

How many calls per District were there ?

```
> nt[,.N,by=District]
```

	District	N
1:	11	21798
2:	1	12107
3:	15	14385
4:	2	13448
5:	3	17649
6:	20	5674
7:	8	22386
8:	5	15258
9:	7	20150
10:	6	19232
11:	18	14178
12:	16	10753
13:	12	8774
14:	4	19789
15:	17	9673
16:	24	9498
17:	9	16656
18:	19	15608
19:	22	10745
20:	14	12537
21:	10	15016
22:	25	19658
23:	13	7084
24:	NA	2079
25:	31	6

Aggregation on a Large File

How many calls per District were there ?

```
nt[, .N, by=District]
```

	District	N
1:	11	21798
2:	1	12107
3:	15	14385
4:	2	13448
5:	3	17649
6:	20	5674
7:	8	22386
8:	5	15258
9:	7	20150
10:	6	19232
11:	18	14178
12:	16	10753
13:	12	8774
14:	4	19789
15:	17	9673
16:	24	9498
17:	9	16656
18:	19	15608
19:	22	10745
20:	14	12537
21:	10	15016
22:	25	19658
23:	13	7084
24:	NA	2079
25:	31	6

Aggregation on a Large File

- How many calls per District were there ?
- Notice that when we use a function designed for a native data frame that it will work
- However, as we will soon see, it works much slower than using the `data.table` approach

```
table(nt$District)
```

1	2	3	4	5	6	7	8	9	10	11	12	13
12107	13448	17649	19789	15258	19232	20150	22386	16656	15016	21798	8774	7084
14	15	16	17	18	19	20	22	24	25	31		
12537	14385	10753	9673	14178	15608	5674	10745	9498	19658	6		

Aggregation on a Large File

The difference in time is pretty significant although from the user point of view it might not be that different. The utility of `data.table` becomes apparent when reading in really large files (millions of rows and lots of attributes/columns)

```
system.time( nt[,.N,by=District] )
```

user	system	elapsed
0.003	0.000	0.002

```
system.time( table(nt$District) )
```

user	system	elapsed
0.089	0.004	0.093

Aggregation on a Large File

Let's randomly sample 500 rows and then find the mean calls to the cops as grouped by FBI.Code (whatever that corresponds to) check

<https://www2.fbi.gov/ucr/nibrs/manuals/v1all.pdf> to see them all

```
nt[sample(1:.N,500),.(mean=mean(.N)),by=FBI.Code]
```

	FBI.Code	mean
1:	08A	28
2:	08B	97
3:	06	98
4:	24	8
5:	03	21
6:	11	13
7:	22	3
8:	26	38
9:	05	33
10:	04A	10
11:	14	51
12:	18	47
13:	07	18
14:	15	15

Wikipedia Page Traffic Statistics

This dataset contains a 150 GB sample of the data used to power trendingtopics.org.

It includes a full 3 months of hourly page traffic statistics from Wikipedia (1/1/2011-3/31/2011)

- Contains hourly wikipedia article traffic statistics dataset covering 3 month period from January 01 2011 to March 31 2011
- Each of the 2,161 log files is named with the date and time of collection: pagecounts-20090430-230000.gz
- Each line has 4 fields: projectcode, pagename, pageviews, bytes

Data available from Amazon Public Data Sets at <https://aws.amazon.com/datasets/wikipedia-page-traffic-statistic-v3/>

Wikipedia Page Traffic Statistics

I took some of these files and combined them into a single file of size 1.4GB and 31,164,567 records with 4 columns. This is a small fraction of the total data.

I am running this test on an Amazon Web Services instance that has 4GB of memory and 2 Cores.

Many laptops have at least this and usually more.

```
$ ls -lh combined_wiki.txt
-rw-r--r-- 1 ubuntu root 1.4G Sep  3 19:27 combined_wiki.txt
```

```
$ wc -l combined_wiki.txt
31164567 combined_wiki.txt
```

Wikipedia Page Traffic Statistics

Let's read this file in using the native **read.csv** function and then with the **fread** function supplied with the data.table package.

```
system.time(df <- read.csv("combined_wiki.txt",sep=" ",header=F))
```

```
      user  system elapsed  
660.344    4.244  665.825
```

```
nrow(df)  
[1] 31164567
```

```
head(df,5)
```

	V1		V2	V3	V4
1	aa.b	Main_Page	1	5565	
2	aa.b	MediaWiki:Image_sample	1	5179	
3	aa.b	MediaWiki:Upload_source_file	1	5195	
4	aa.b	Wikibooks:Privacy_policy	1	4925	
5	aa.d	MediaWiki:Group-abusefilter-member	1	4912	

Wikipedia Page Traffic Statistics

Let's read this file in using the native **read.csv** function and then with the **fread** function supplied with the data.table package.

```
system.time(dt <- fread("combined_wiki.txt"))
```

Read 31164567 rows and 4 (of 4) columns from 1.307 GB file in 00:02:22

```
      user  system elapsed  
140.476    1.064  156.104
```

```
dt[1:5,]
```

	V1	V2	V3	V4
1:	aa.b	Main_Page	1	5565
2:	aa.b	MediaWiki:Image_sample	1	5179
3:	aa.b	MediaWiki:Upload_source_file	1	5195
4:	aa.b	Wikibooks:Privacy_policy	1	4925
5:	aa.d	MediaWiki:Group-abusefilter-member	1	4912

```
dt[,.N]
```

```
[1] 31164567
```

Wikipedia Page Traffic Statistics

For all rows where page access is > 5 what is the mean number of page accesses as grouped by project code ?

```
> dt[V3 > 5,.(mean=mean(V3)),by=V1]
```

	V1	mean
1:	ab	8.777778
2:	ace	9.785714
3:	af	27.277778
4:	als	13.964286
5:	am	10.900000

597:	zh-min-nan.d	6.000000
598:	zu.b	6.000000
599:	ms.mw	6.000000
600:	mt.d	6.000000
601:	tg.d	10.000000

Wikipedia Page Traffic Statistics

For all rows where page access is > 5 what is the mean number of page accesses as grouped by project code ?

```
system.time( dt[V3 > 5,.(mean=mean(V3)),by=V1] )
      user  system elapsed
0.412    0.080    0.492
```

setkey function

It is possible to sort a data table by using the **setkey(DT, key)** function that is part of the **data.table** package

This will reorganize the data table DT by key where key is a column name

This isn't necessary unless perhaps you want to pre sort the data table for some aggregation in which case pre sorted the data can help.

to make this clear let's look at the smaller data frame from the first slides

```
DT <- data.table(x=c("B","A","B","A","B"),y=c(7,2,1,5,9))
```

	x	y
1:	B	7
2:	A	2
3:	B	1
4:	A	5
5:	B	9

setkey function

```
setkey(DT,x) # Will sort the table by x
```

	x	y
1:	A	2
2:	A	5
3:	B	7
4:	B	1
5:	B	9

```
setkey(DT,y) # Will sort the table by y
```

	x	y
1:	B	1
2:	A	2
3:	A	5
4:	B	7
5:	B	9

setkey function

```
setkey(DT,x,y)  # Sort by x and y
```

DT

	x	y
1:	A	2
2:	A	5
3:	B	1
4:	B	7
5:	B	9

Review

Why did I show you all this ? Let's review the good stuff about `data.table`

- Excellent for reading in very large data sets fast
- Provides a way to do aggregation within the brackets
- Can behave just like a data frame if need be
- It is easy to turn an existing data frame into a data table

Okay why wouldn't you use `data.table` ?

- Do you REALLY need to read in a 100 million row dataset ?
- Just read in a fraction of the data
- The bracket notation can get busy
- Maybe you should use a data base / SQL approach

sqldf

There is another way to do this. Turn the data frame into a SQLite database We will explore SQL and databses in another class

```
library(sqldf)

sqldf("attach cwiki as new")      # create a new SQLite database

# Read in the .txt file into a table

system.time( read.csv.sql("combined_wiki.txt",
  sql="create table info as select * from file",
  dbname="cwiki",header=FALSE,sep=" ")
)

user    system elapsed
85.832  10.676 134.189
```

sqldf

Prove that we got all the records

```
sqldf("select count(*) from info", dbname="cwiki")  
      count(*)  
1 31164567
```

How many distinct Project Codes are there ?

```
sqldf("select count(distinct(V1)) from info",dbname="cwiki")  
      count(distinct(V1))  
1                1266
```

How many records correspond to project code aa.b ?

```
sqldf("select count(*) from info where V1 = 'aa.b' ",  
      dbname="cwiki")  
      count(*)  
1          52
```

sqldf

```
# Create a data frame holding all rows relating to aa.b
```

```
df <- sqldf("select * from info where V1 = 'aa.b' ",  
            dbname="cwiki")
```

```
head(df)
```

	V1	V2	V3	V4
1	aa.b	Main_Page	1	5565
2	aa.b	MediaWiki:Image_sample	1	5179
3	aa.b	MediaWiki:Upload_source_file	1	5195
4	aa.b	Wikibooks:Privacy_policy	1	4925
5	aa.b	Main_Page	3	27714
6	aa.b	Special:Imagelist	1	540

sqldf

```
# Compute the average number of megabytes downloaded for  
# the top 5 unique project codes
```

```
sqldf("select distinct(V1) as ProjCode,avg(V4)/1000000 as MB  
      from info group by V1 order by MB  
      desc limit 5",dbname="cwiki")
```

	ProjCode	MB
1	en.mw	77518.224
2	ja.mw	9126.983
3	fr.mw	2020.454
4	ru.mw	1311.165
5	de.mw	1214.592

sqldf

Well if you knew that you just wanted the Project codes aa.b from the beginning then you could have specified this when creating the database originally.

This would save some time in the long run since we wouldn't be reading in data for other project codes

```
sqldf("attach cwiki as new")
```

```
# We just select the aa.b project codes
```

```
read.csv.sql("combined_wiki.txt",  
             sql = "create table info as select * from file  
                   where V1 = 'aa.b'",  
             dbname = "cwiki", header = FALSE, sep = " ")
```

```
sqldf("select count(*) from info",dbname="cwiki")  
count(*)
```

```
1      52
```